

Lab 2

Setting up a Cluster

Objectives:

1. Learn basic hardware and software setup for a simple cluster.
2. Learn basic message passing programming environment.

In this lab, you will learn how to set up a parallel computer cluster with two Dell nodes using the Ubuntu operating system. Ubuntu is a popular Linux distribution mostly used for personal computer / laptop, but has steadily gain user base in the server sector.

As a practice, you will add a Nvidia Jetson TK1 node to the cluster in the later part of the exercise.

As there are only 7 sets of hardware components (i.e. enough for 7 clusters), you have to form a team of **3 members** for this lab.

A. Setting up a 2-node cluster


There are a total of 5 steps to setup the 2-node cluster:

1. Set up the hardware
2. Install Ubuntu
3. Install essential software
4. Configure SSH access
5. Test cluster by running an MPI¹ parallel program

Step 1: Set up Hardware

To setup a parallel cluster consisting of two nodes, connect the hardware in your workbench as follows:

¹ MPI (Message-passing library interface) is a parallel programming library for writing message-passing parallel program on distributed memory system. We will cover MPI later in the lecture. In this lab you will only run a simple hello world program to check the communication among processes running on different processors (cores) in the parallel cluster system.



Node 1 (8 cores) is a **Dell Optiplex 990** with tower form factor and node 2 (4 cores) is a **Dell Inspiron** all-in-one desktop (see Appendix for photos).

Step 2: Install Ubuntu

Install Ubuntu 14.04 Server LTS² on all nodes using the given Live CD. Detailed screenshots can be found at <http://ubuntuserverguide.com/2014/04/how-to-install-ubuntu-server-14-04-trusty-tahr.html>

The main steps to install Ubuntu Server 14.04 LTS are (trivial steps are omitted!):

- “Select language”, “Select Install Ubuntu Server”, “Select installation language”, “Select your location”, “Configure the keyboard” ...
- Next, “Configure the network”, choose the Ethernet interface, then fill in the hostname of your node (the hostname is on the yellow label on the monitors, e.g., soctf-pdc-002 and soctf-pdc-009).
- “Set up users and passwords” – set up with the following user on each node:

	Node 1 (Dell Optiplex)	Node 2 (Dell Inspiron)
Full name	node1	node2
Username	node1	node2
Password	node1	node2

When prompted to "encrypt your home directory": *No*

² The latest stable built is 16.04 as of August, 2016. However, to ensure compatibility with all different platforms, we have opted to use a slightly older version.

- d. Next, “Partition disks”. When prompted to "unmount partitions that are in use", select **"Yes"**.

Choose “**Manual**” partitioning as the partitioning method. From the list of partitions, select the partition with the name ***FREE SPACE***.

Note: The list of partitions on your machine may be different. Please consult the tutor.

Next, select “*Automatically partition the free space*”, select “*Finish partitioning and write changes to disk*”.

The tutor will inspect your setup before you proceed to the next step. Changes to the harddisk partition is hard to reverse, so please be patient and verify before proceed.

After your setup has been inspected, choose “*Write changes to disks: Yes*”.

- e. Configure HTTP Proxy: Leave blank and select ***Continue***
- f. Configure Taskel: select “*No automatic updates*”
- g. Software selection.

Do not press <Enter>. Use <Space Bar> to select the necessary software.

Select “*OpenSSH server*” then ***Continue***

- h. Install the GRUB boot loader on a hard disk: ***Yes***

If the installation is successful, “Installation complete” is displayed. Select “Continue” to reboot the system. Next, login to each node using the username and password as defined earlier.

Step 3: Install Essential Software

Install the package “build-essential” (include make, gcc, etc. software development tools) on all nodes:

```
$ sudo apt-get install build-essential
```

If you are prompted for password, just enter the password of the user (e.g. User: "Node1" → Password: "Node1").

Install “OpenMPI”, an implementation of MPI, on all nodes.

```
$ sudo apt-get install libopenmpi-dev openmpi-bin openmpi-doc
```

Install the package “vim” editor on all nodes:

```
$ sudo apt-get install vim
```

Installing “OpenSSH” client and server on all nodes.

```
$ sudo apt-get install openssh-server openssh-client
```

Step 4: Configure SSH Access

The nodes will communicate each other using their full hostname. For example, the full hostname of a node with hostname **soctf-pdc-002** is **soctf-pdc-002.comp.nus.edu.sg**.

Assume we have the following two nodes:

Node	Full Hostname
node1	soctf-pdc- <u>002</u> .comp.nus.edu.sg
node2	soctf-pdc- <u>009</u> .comp.nus.edu.sg

First, ensure the node can reach each other via network.

On node **node1**:

```
$ ping soctf-pdc-009.comp.nus.edu.sg
```

On node **node2**:

```
$ ping soctf-pdc-002.comp.nus.edu.sg
```

On both nodes, create a new user called **cs3210**:

```
$ sudo useradd -s /bin/bash -d /home/cs3210 -m cs3210
```

```
$ sudo passwd cs3210
```

Enter “pc2016” as the password.

Exit from the current user.

```
$ exit
```

To enable SSH authentication without password, we need to generate and exchange public/private key pairs on both nodes:

- i. login with username **cs3210** and password **pc2016**.
- ii. Generate public/private key pair.

```
$ssh-keygen -t dsa
```

[Press enter a few times to ignore the prompts]

This command generates two files: `id_dsa` and `id_dsa.pub` in the `~/.ssh` folder.

`id_dsa` is the private key and `id_dsa.pub` is the public key.

We need to keep the private key and copy the public key to the other node.

To exchange the public keys, use the **ssh-copy-id** command to copy `id_dsa.pub`. Also, to facilitate MPI, we will also add **the node's own public key** to the authorized list.

On node **node1** (copy its public key, `id_dsa.pub`, to node2):

```
$cd ~/.ssh
```

```
$ssh-copy-id -i id_dsa.pub soctf-pdc-009.comp.nus.edu.sg
```

```
$cat id_dsa.pub >> authorized_keys
```

On node **node2** (copy its `id_dsa.pub` to node1):

```
$cd ~/.ssh
```

```
$ ssh-copy-id -i id_dsa.pub soctf-pdc-002.comp.nus.edu.sg
```

```
$cat id_dsa.pub >> authorized_keys
```

Now, node1 and node2 should be able to SSH each other without password. Do the following check on each node (change the hostname accordingly):

```
$ ssh cs3210@soctf-pdc-009.comp.nus.edu.sg
```

[Check that you can login without password]

```
$ exit
```

Caution: As we have the same user id on both nodes, it may be easily confused (for example, SSH into **node2** and mistakenly thought that you are still working with **node1** etc). So, refrain from SSH into another node after this check.

Step 5: Test cluster by running an MPI parallel program

In this step, we are going to run an MPI program on the cluster. On **node1**, download `hello.c` program:

```
$ wget www.comp.nus.edu.sg/~sooyj/cs3210/Lab2/hello.c
```

You can find the full source code in the appendix.

First, compile the program on **node1**:

```
$ mpicc hello.c -o hello
```

Next, copy the binary to node **node2**.

```
$ scp hello cs3210@soctf-pdc-009.comp.nus.edu.sg:~
```

In order to run the MPI program on the cluster, create a file called `hostfile` using the editor `vim` (or any text editor you are familiar with) as follows:

```
$vim hostfile
```

(to start editing in vim, press <Esc>:i)

Add the list of hostnames of nodes in the cluster as follows:

```
soctf-pdc-002.comp.nus.edu.sg
```

```
soctf-pdc-009.comp.nus.edu.sg
```

(to save and exit from vim press <Esc>:wq)

Lastly, execute the program:

```
$mpirun -hostfile hostfile -np 2 hello
```

```
Hello world from process 0 of 2 running on soctf-pdc-002
```

```
Hello world from process 1 of 2 running on soctf-pdc-009
```

Well done, you have successfully setup a working parallel system!

You can try to increase the number of independent process from 2 to 4:

```
$mpirun -hostfile hostfile -np 4 hello
```

to observe the output.

B. Adding a Nvidia Jetson TK1 node to the cluster

Ubuntu 14.04 LTS is already installed in the embedded-class low-power Jetson TK1³ system and a user account, *cs3210*, has been created. Your tasks are:

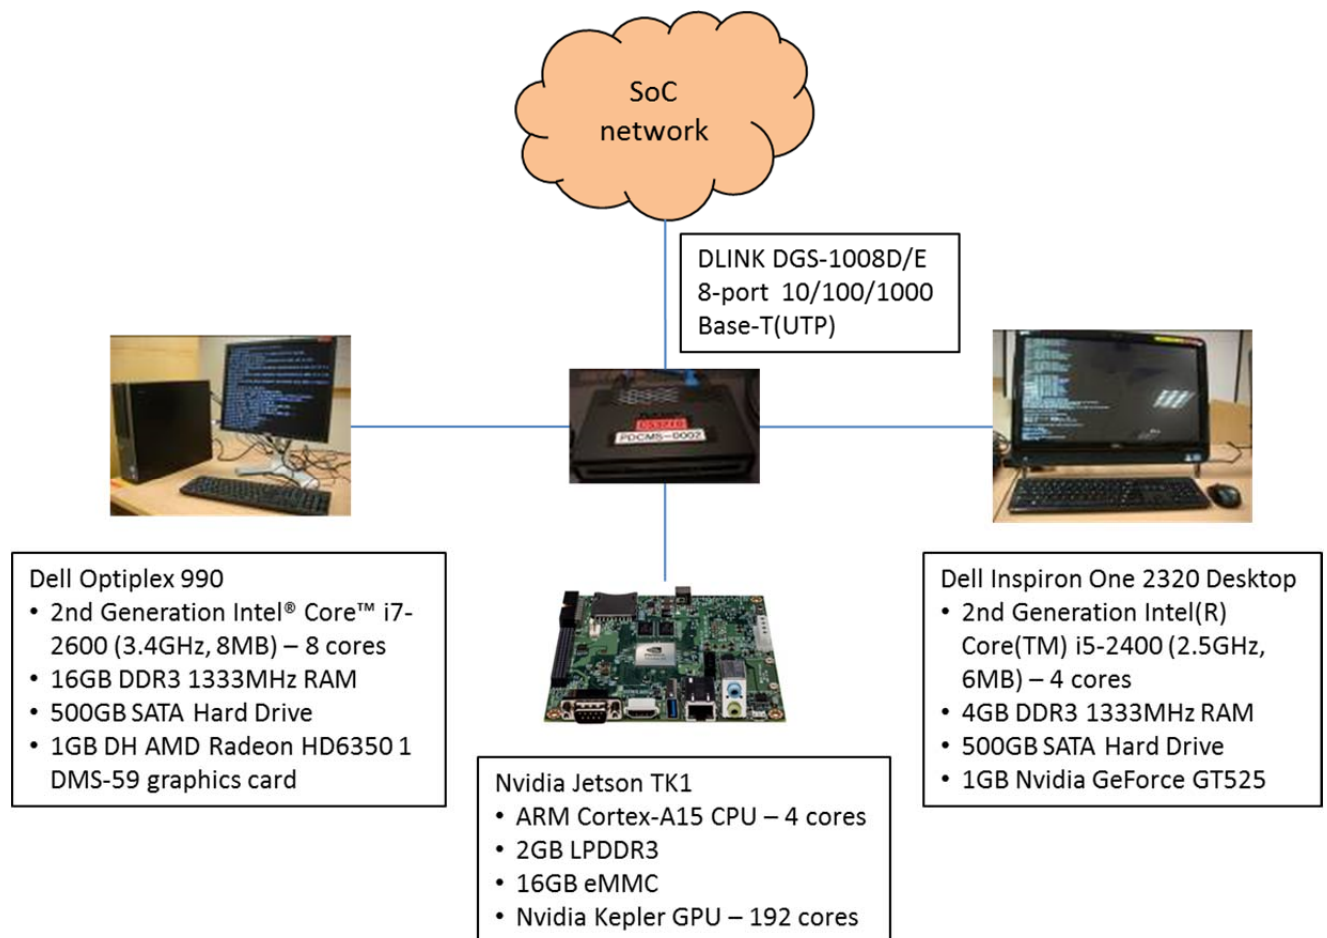
1. Add this system to the 2-node cluster
2. Configure the ssh access (same step 4 in section A)
3. Run the MPI “Hello World” program on the 3-node cluster (similar to step 5 in section A). (Hint: node1 and node2 are both x86 machines, while jetson is on ARM processor).

Once your 3-node cluster is working, please ask the TA to verify.

³ More about Jetson TK1 at <http://www.nvidia.com/object/jetson-tk1-embedded-dev-kit.html>

Appendix

Cluster Configuration



Installing Network Hardware

Plug one end of the cable into a node and the other end into a switch



Source Code for the MPI Program `hello.c`

```
/* C Example */
#include <stdio.h>
#include <mpi.h>

int main (int argc, char *argv[])
{
    int rank, size;
    char msg[80];
    int length;

    /* starts MPI */
    MPI_Init (&argc, &argv);

    /* get current process id */
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);

    /* get number of processes */
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    /* Gets the name of the processor (computer name) */
    MPI_Get_processor_name(msg, &length);

    printf( "Hello world from process %d of %d running on %s
\n", rank, size, msg);

    /* exiting */
    MPI_Finalize();
    return 0;
}
```